



The One Billion Row Challenge:

je PHP rychlejší než Java?

27.11.2024 v 18:00

Restaurace U Salzmannů

Salonek v 1. patře

Vstup a občerstvení zdarma



PeoplePath

Few facts about me...



- Software architect
- +25 years experience in IT.
- Contributor to PHP magazines.
- Conference speaker.

Few facts about me...



- Demotivational speaker.
- Loves crazy projects.
- Tried to patent QR payments in the USA.

Languages that I
use the most...



Code

Costumes

Sounds

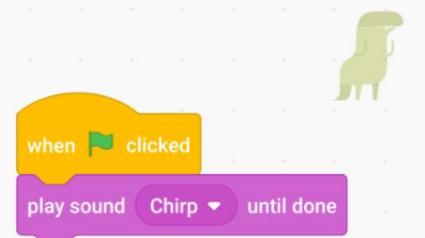


Sound

play sound Chirp until done

start sound Chirp

stop all sounds



SCRATCH

Sprite Sprite1 x 0 y 0
 Size 100 Direction 90

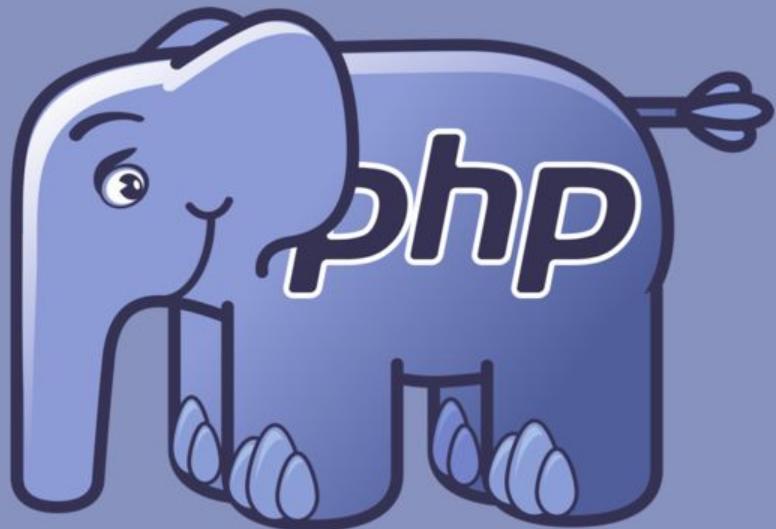


Stage



Backdrops

1





Error 502

Bad Gateway



BROWSER
Working



MYRA
Working



HOST
Error

What happened?

The communication with the origin web server timed out or was invalid. Therefore the web page cannot be displayed.

What can I do?

If you are a visitor of this website:

Please try again later.

If you are the owner of this website:

Please make sure that your origin server is not under heavy load. Sometimes a crashing program can also be the root cause.

opcache.jit_debug int

A bit mask specifying which JIT debug output to enable. [For possible values, please consult » zend_jit.h](#) (search for macro definitions beginning with ZEND_JIT_DEBUG).

opcache.jit_bisect_limit int

Debugging option that disables JIT compilation after compiling a certain number of functions. This may be helpful to bisect the source of a JIT miscompilation. Note: this option only works when JIT trigger is set to 0 (compile on script load) or 1 (compile on first execution), e.g. opcache.jit=1215. See more in [opcache.jit](#) option.

opcache.jit_prof_threshold float

When using the "profile on first request" trigger mode, this threshold determines whether a function is considered hot. The number of calls to the function divided by the number of calls to all functions must be above the threshold. For example, a threshold of 0.005 means that functions that made up more than 0.5% of all calls will be JIT compiled.

opcache.jit_max_root_traces int

Maximum number of root traces. The root trace is an execution flow taking one path through the code firstly, which is a unit of JIT compilation. JIT will not compile new code if it reaches this limit.

opcache.jit_max_side_traces int

Maximum number of side traces a root trace may have. The side trace is another execution flow that does not follow the path of compiled root trace. Side traces belonging to the same root trace will not be compiled if it reaches this limit.

opcache.jit_max_exit_counters int

Maximum number of side trace exit counters. This limits the total number of side traces there may be, across all root traces.

opcache.jit_hot_loop int

After how many iterations a loop is considered hot. Valid value range is [0,255]; for any setting out of this range, e.g. -1 or 256, default value will be used instead. 0 will disable JIT to trace and compile any loops.

opcache.jit_hot_func int

After how many calls a function is considered hot. Valid value range is [0,255]; for any setting out of this range, e.g. -1 or 256, default value will be used instead. 0 will disable JIT to trace and compile any functions.

opcache.jit_hot_return int

After how many returns a return is considered hot. Valid value range is [0,255]; for any setting out of this range, e.g. -1 or 256, default value will be used instead. 0 will disable JIT to trace and compile any returns.

opcache.jit_hot_side_exit int

After how many exits a side exit is considered hot. Valid value range is [0,255]; for any setting out of this range, e.g. -1 or 256, default



**YOU'RE NOT A
BAD PROGRAMMER**

**IT'S A BAD
PROGRAMMING LANGUAGE**

1BRC



The One Billion Row Challenge

1BRC - objective

- Calculate stats from 1B-row file.

Hamburg;12.0

Bulawayo;8.9

Palembang;38.8

Hamburg;5.0

London;0.8

1BRC - output

- Print min, mean, and max temps per station alphabetically.

{Abha=5.0/18.0/27.4, Abidjan=15.7/26.0/34.1, ...}

1BRC - background

- Originally in JAVA.
- Other languages C, Python, RUST, etc.
- In PHP done by Florian Engelhardt.

TL;DR - funny challenge for optimisation



Measure methodology 1 BRC

- Using wall clock, on Linux.
- Results differ across hardware.
- Flushing vmda caches after each run.
- CLI tool hyperfine



“Talk is
cheap. Show
me the code.”

Linus Torvalds

```
1 <?php
2 $stations = []; // create empty array
3 $fp = fopen('measurements.txt', 'r'); // open file for read
4 while ($data = fgetcsv($fp, null, ',')) { // get line from file and parse it
5     if (!isset($stations[$data[0]])) {
6         $stations[$data[0]] = [
7             $data[1], // minimum value
8             $data[1], // maximum value
9             $data[1], // sum of values
10            1 // count of measurements
11        ];
12    } else {
13        $stations[$data[0]][3]++; // Increment the count
14        $stations[$data[0]][2] += $data[1]; // Add to the sum
15        if ($data[1] < $stations[$data[0]][0]) {
16            $stations[$data[0]][0] = $data[1]; // Update the minimum value
17        }
18        if ($data[1] > $stations[$data[0]][1]) {
19            $stations[$data[0]][1] = $data[1]; // Update the maximum value
20        }
21    }
22 }
23 ksort($stations); // Sort an array by key in ascending order
24 echo '{';
25 foreach($stations as $k=>&$station) {
26     // Calculate the average value
27     $station[2] = $station[2]/$station[3];
28     echo $k, '=', $station[0], '/', $station[2], '/', $station[1], ", \n";
29 }
30 echo '}';
```



```
1 <?php  
2 $stations = [];  
3 $fp = fopen('measurements.txt', 'r');  
4 while ($data = fgetcsv($fp, null, ';')) {
```

- Array is an ordered map (KV, **hash table**).
- fgetcsv - get line from file and parse it.
 - Eg.: berlin;5.87 will return array of two elements: berlin and 5.87

```
4 while ($data = fgetcsv($fp, null, ';')) {  
5     if (!isset($stations[$data[0]])) {  
6         $stations[$data[0]] = [  
7             $data[1], // minimum value  
8             $data[1], // maximum value  
9             $data[1], // sum of values  
10            1           // count of measurements  
11        ];  
12    } else {
```

```
12 } else {
13     $stations[$data[0]][3]++;
14     $stations[$data[0]][2] += $data[1];
15     if ($data[1] < $stations[$data[0]][0]) {
16         $stations[$data[0]][0] = $data[1];
17     }
18     if ($data[1] > $stations[$data[0]][1]) {
19         $stations[$data[0]][1] = $data[1];
20     }
21 }
22 } // while Loop end
```

```
22 } // while Loop end
23 ksort($stations);
24 echo '{';
25 foreach($stations as $k=>&$station) {
26     // Calculate the average value
27     $station[2] = $station[2]/$station[3];
28     echo $k, '=', $station[0], '/', $station[2], '/',
29         $station[1], ", \n";
30 }
```

What are arrays in PHP?

- A. A collection of random values stored in memory.
- B. A mystical portal that stores data in parallel universes.
- C. An ordered map, which is a type of hash table.
- D. A fixed-size collection of integers.

What are arrays in PHP?

- A collection of random values stored in memory.



- A mystical portal that stores data in parallel universes.



- An ordered map, which is a type of hash table.

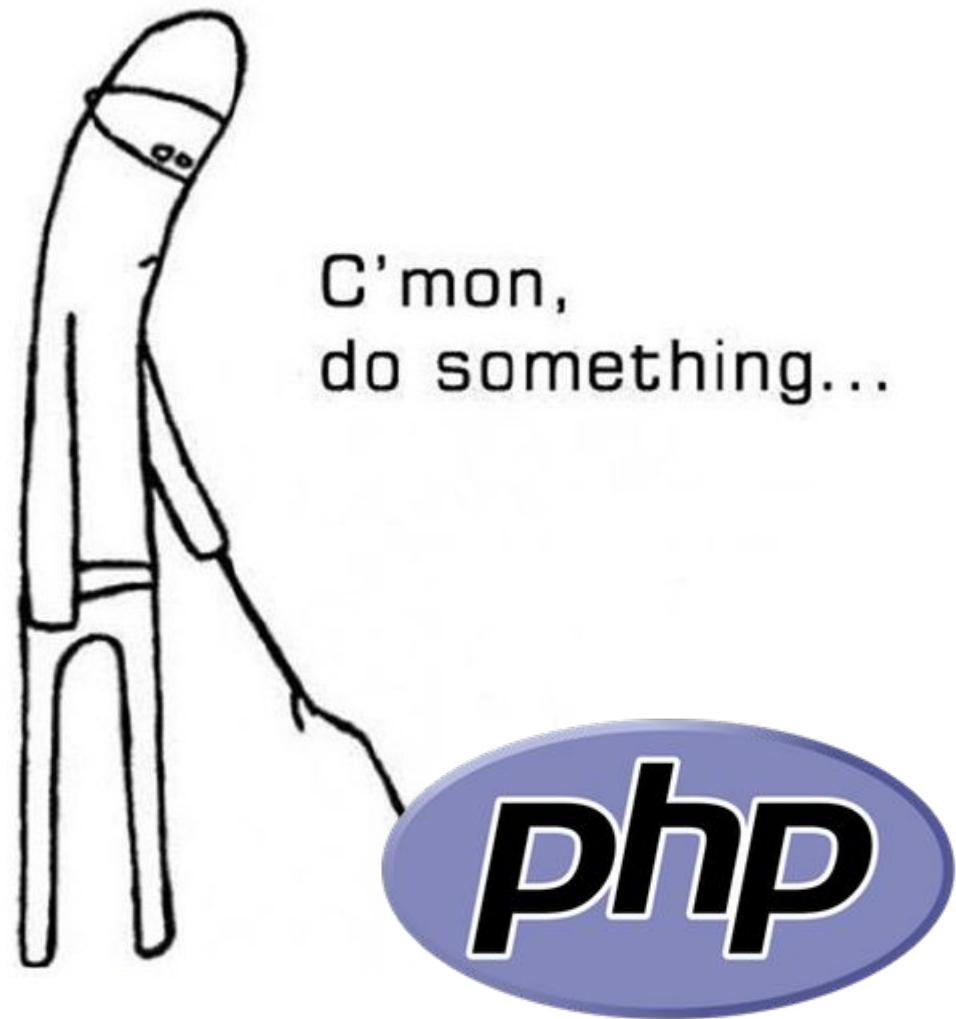


- A fixed-size collection of integers.



Baseline

902.70 sec



C'mon,
do something...

What to do?

What to do?

- A. Use other language?
- B. Consult with more experience developer?
- C. Don't change anything, check configuration?
- D. Act like nothing is happening!

What to do?

- Use other language?



- Consult with more experience developer?



- Don't change anything, check configuration?



- Act like nothing is happening!



**Don't change
anything, check
configuration!**

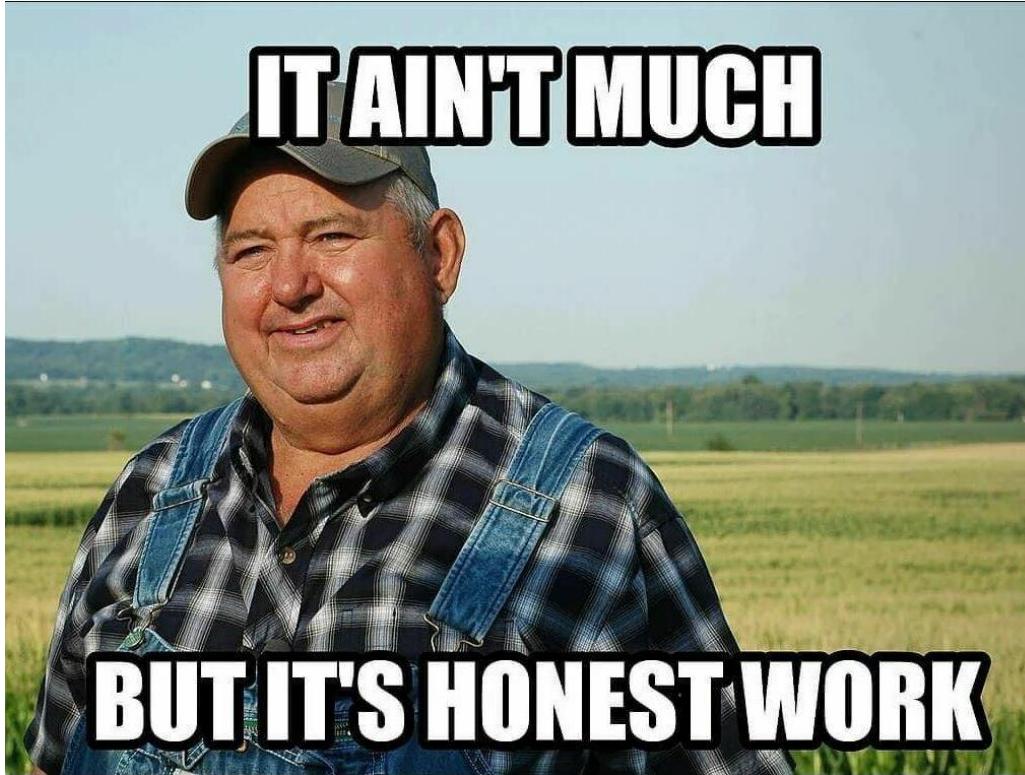
Don't change anything

Version	Time (sec)	Raw difference	% improvement
Baseline	902.70	n/a	n/a
OpCache	882.97	19.73	2.19%
JIT	850.38	32.59	3.69%

- 5.8% (52.32 sec) just by configuration.

IT AINT MUCH

BUT IT'S HONEST WORK



OpCache

- Can be enabled also for CLI.
- Up to 6 optimisation passes.
- JIT engine.

**TL;DR Check
configuration and
runtime**

What's next?

What's next?

- A. Randomly changing things?
- B. Do the profiling?
- C. C'mon 850 sec? It's fast enough!
- D. Gave up programing, start farming?

What's next?

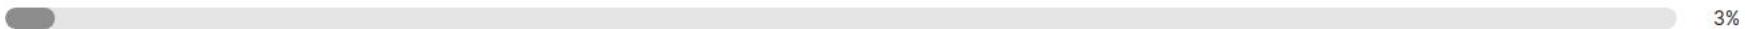
Randomly changing things?



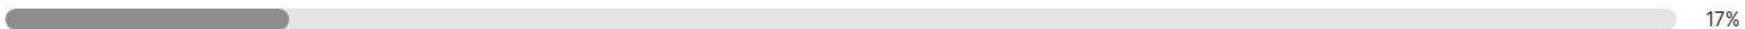
Do the profiling?



C'mon 850 sec? It's fast enough!



Gave up programing, start farming?



C'mon 850 sec? It's fast enough!



Croak NO! Fix it, croak!

Profiling

Profiling - PHP specific

- Xdebug (tracing profiler)
- Reli-prof (sampling profiler)
- Tideways (SASS)

XDebug

Open... | < Back | > Forward | ^ Up | % Relative | Cycle Detection | Relative to Parent | Shorten Templates | Time_(10ns) | ▾

Flat Profile

Search: Search Query (No Grouping)

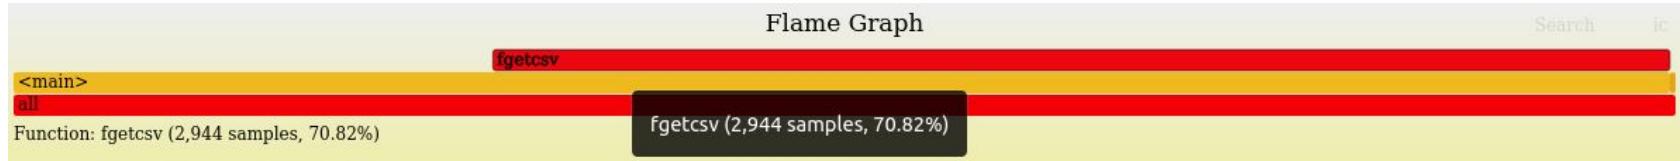
{main}

Incl.	Self	Called	Function	Loc
100.00	63.50	(0)	{main}	0-n
35.60	35.60	10 000 001	php:fgetcsv	php
0.90	0.90	1	php:fopen	php
0.00	0.00	1	php:ksort	php

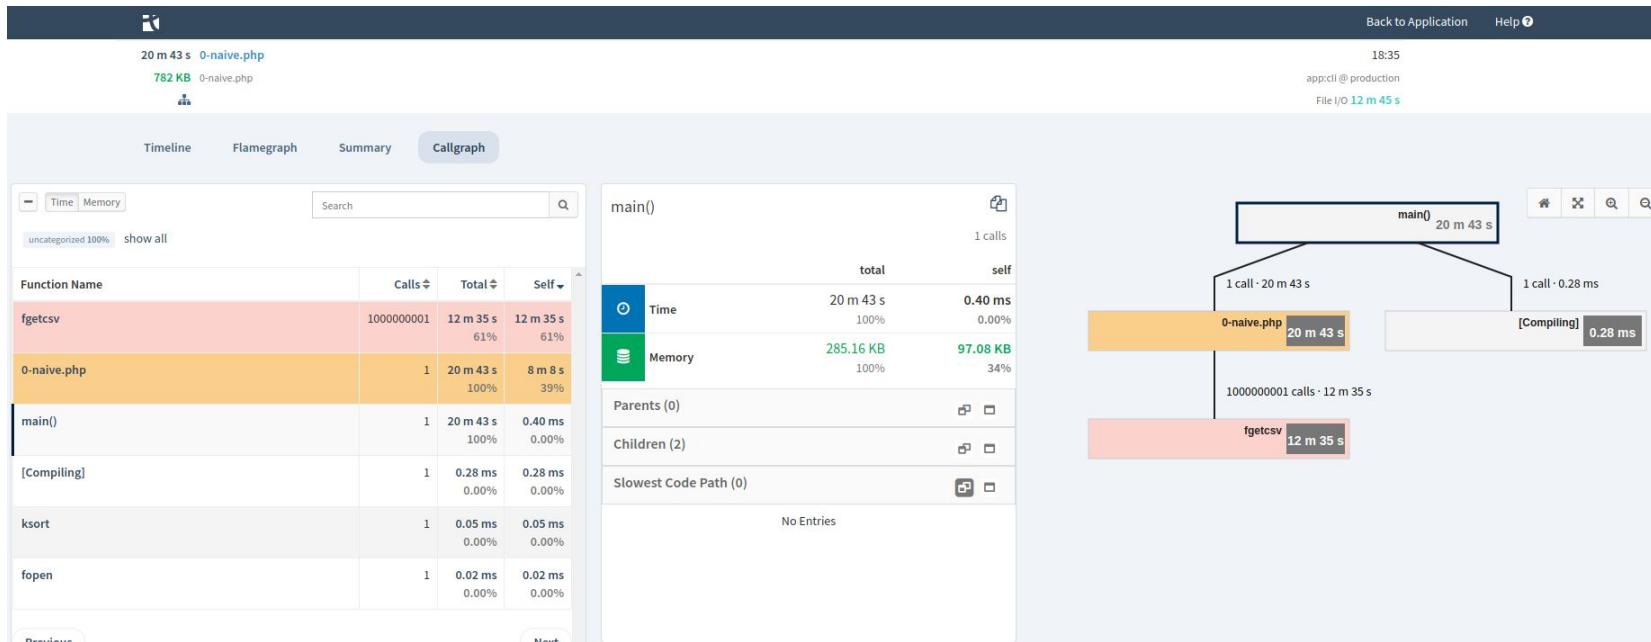
Types Callers All Callers Callee Map Source Code

#	Time_(10ns)	Time_(10ns)	Source
0			--- From '/home/bohuslavsimek/Public/TEST-1brc/0-naive.php' ---
1	63.50	63.50	<?php
2			\$stations = [];
3			\$fp = fopen('measurements-medium3.txt', 'r');
4	0.90	0.90	■ 1 call to 'php:fopen' (php:internal)
5			while (\$data = fgetcsv(\$fp, null, ',')) {
6	35.60	35.60	■ 10000001 calls to 'php:fgetcsv' (php:internal)
7			if (!isset(\$stations[\$data[0]])) {
8			\$stations[\$data[0]] = [
9			\$data[1],
10			...
11			}
12			...
13			}
14			ksort(\$stations);
15		0.00	■ 1 call to 'php:ksort' (php:internal)
16			echo '{';
17			foreach(\$stations as \$k=>&\$station) {

Reli-prof (sampling profiler)



Tideways (SASS)



Profiling - universal tools

- Spread time function - “shotgun” method.
- Blackfire (SASS)
- Datadog (SASS)
- Perf (Linux)

2: sudo perf report

Samples: 3K of event 'cycles:P', Event count (approx.): 145919026547

Overhead	Com	Shared Object	Symbol
31,62%	php	libc.so.6	[.] gconv_transform_utf8_internal
21,26%	php	libc.so.6	[.] __mbrtowc
6,27%	php	php	[.] zend strtod
5,66%	php	php	[.] php_fgetcsv_lookup_trailing_spaces
4,21%	php	php	[.] __is_numeric_string_ex.part.0
4,05%	php	php	[.] php_fgetcsv
2,62%	php	libc.so.6	[.] __dl_mcount_wrapper_check
1,84%	php	php	[.] zend_i smart_strcmp
1,53%	php	php	[.] __php_stream_get_line
1,46%	php	php	[.] zend_hash_find
1,15%	php	php	[.] zif_fgetcsv
0,94%	php	php	[.] add_function_slow
0,90%	php	php	[.] zend_string_equal_val
0,89%	php	php	[.] zend_hash_next_index_insert
0,85%	php	php	[.] __efree
0,76%	php	libc.so.6	[.] __memmove_avx_unaligned_erms
0,76%	php	php	[.] zend_compare
0,72%	php	php	[.] __emalloc
0,64%	php	php	[.] __mbrtowc@plt
0,61%	php	opcache.so	[.] zend_jit_symtable_lookup_rw
0,60%	php	php	[.] zend_hash_func
0,54%	php	php	[.] ZEND_PRE_INC_SPEC_VAR_RETVAL_UNUSED_HANDLER
0,53%	php	php	[.] zend_i try_convert_scalar_to_number.isra.0
0,45%	php	php	[.] add_function
0,40%	php	php	[.] zend_array_destroy
0,33%	php	php	[.] php_stream_locate_eol
0,28%	php	php	[.] zend_mm_alloc
0,28%	php	php	[.] zend_new_array_0
0,22%	php	php	[.] zend_fetch_resource2
0,22%	php	libc.so.6	[.] __ctype_b_loc
0,22%	php	libc.so.6	[.] __memchr_avx2
0,20%	php	php	[.] zend_is_true
0,20%	php	php	[.] __emalloc_160
0,20%	php	php	[.] __realloc
0,11%	php	php	[.] zval_ptr_dtor
0,11%	php	php	[.] memcpy@plt
0,11%	php	php	[.] __emalloc_56
0,09%	php	[JIT] tid 31918	[.] 0x00005555c4600bf5

3: sudo perf report

Samples: 3K of event 'cycles:P', Event count (approx.): 145929078431, Thread: php, DSO: php

Children	Self	Com	Symbol
+ 13,97%	5,98%	php	[.] php_fgetcsv_lookup_trailing_spaces
+ 12,67%	5,87%	php	[.] zend strtod
+ 8,96%	4,36%	php	[.] php_fgetcsv
+ 6,22%	4,08%	php	[.] _is_numeric_string_ex.part.0
+ 5,56%	0,00%	php	[.] zend_execute
+ 4,87%	1,60%	php	[.] zend_hash_find
+ 3,57%	1,61%	php	[.] zend_smart_strcmp
+ 2,81%	0,00%	php	[.] php_stdio_read
+ 2,44%	1,51%	php	[.] _php_stream_get_line
+ 2,44%	0,91%	php	[.] _efree
+ 2,28%	0,37%	php	[.] _emalloc
+ 1,95%	1,27%	php	[.] zend_hash_next_index_insert
+ 1,88%	1,00%	php	[.] zif_fgetcsv
+ 1,51%	0,77%	php	[.] zend_string_equal_val
+ 1,43%	0,50%	php	[.] zend_compare
+ 1,28%	0,77%	php	[.] add_function_slow
1,12%	0,66%	php	[.] zend_array_destroy
1,02%	0,99%	php	[.] mbrtowc@plt
0,93%	0,54%	php	[.] php_stream_locate_eol
0,86%	0,46%	php	[.] zend_hash_func
0,79%	0,37%	php	[.] _zend_try_convert_scalar_to_number.isra.0
0,65%	0,36%	php	[.] ZEND_PRE_INC_SPEC_VAR_RETVAL_UNUSED_HANDLER
0,62%	0,20%	php	[.] _emalloc_160
0,60%	0,40%	php	[.] _zend_new_array_0
0,56%	0,39%	php	[.] add_function
0,49%	0,26%	php	[.] _erealloc
0,40%	0,17%	php	[.] _emalloc_56
0,37%	0,12%	php	[.] _zend_mm_alloc
0,26%	0,06%	php	[.] _efree_56
0,23%	0,11%	php	[.] zend_is_true
0,23%	0,09%	php	[.] zend_fetch_resource2
0,20%	0,09%	php	[.] rc_dtor_func
0,15%	0,09%	php	[.] zval_ptr_dtor
0,14%	0,08%	php	[.] memcpy@plt
0,12%	0,12%	php	[.] __ctype_b_loc@plt
0,09%	0,03%	php	[.] zend_string_hash_func
0,08%	0,05%	php	[.] php_file_le_stream
0,03%	0,03%	php	[.] php_file_le_pstream

add_function_slow



```
+      1,28%      0,77%  php  [ . ] add_function_slow
```

Perf - slowest functions

1. php_fgetcsv_lookup_trailing_spaces
2. zend strtod
3. php_fgetcsv
4. is_numeric_string_ex.part.0
5. zend_execute
6. zend_hash_find

**TL;DR Find Out
what is slow.**

How to fix this?

- A. Use a higher PHP version to improve performance.
- B. Compress the CSV file to reduce size before reading.
- C. Rewrite the fgetcsv functionality.
- D. Increase the memory limit for the PHP.

How to fix this?

- ⓘ Use a higher PHP version to improve performance.



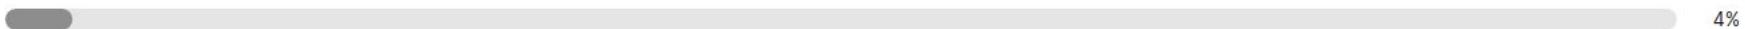
- ⓘ Compress the CSV file to reduce size before reading.



- ✓ Rewrite the fgetcsv functionality.



- ⓘ Increase the memory limit for the PHP.



Replace fgetcsv with fgets



fgets

(PHP 4, PHP 5, PHP 7, PHP 8)

fgets — Gets line from file pointer

Description

```
fgets(resource $stream, ?int $length = null): string|false
```

Gets a line from file pointer.

Parameters

stream

The file pointer must be valid, and must point to a file successfully opened by `fopen()` or `fsockopen()` (and not yet closed by `fclose()`).

Replace fgetcsv with fgets

```
- while ($data = fgetcsv($fp, null, ';')) {
-     if (!isset($stations[$data[0]])) {
-         $stations[$data[0]] = [
+ while ($data = fgets($fp, null)) {
+     $pos = strpos($data, ';');
+     $city = substr($data, 0, $pos);
+     $temp = substr($data, $pos+1, -1);
+     if (!isset($stations[$city])) {
+         $stations[$city] = [
```

Replace `fgetcsv` with `fgets`

Before
850.38 sec



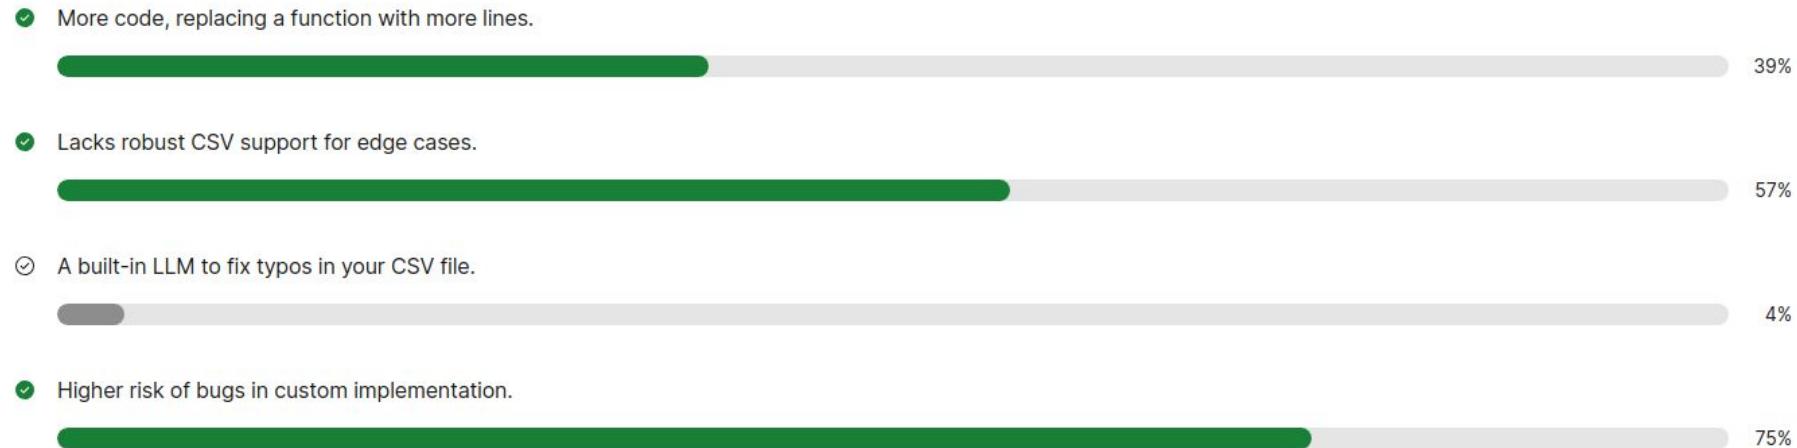
After
294.80 sec

65.33% faster!

Any downside to this replacement?

- A. More code, replacing a function with more lines.
- B. Lacks robust CSV support for edge cases.
- C. A built-in LLM to fix typos in your CSV file.
- D. Higher risk of bugs in custom implementation.

Any downside to this replacement?



TL;DR

**Find out what's
essential.**

TL;DR

**Do only what's
essential.**

**What's next?
Rinse and repeat!**

XDebug

Open... | < Back | > Forward | ^ Up | % Relative | Cycle Detection | Relative to Parent | Shorten Templates | Time_(10ns) | ▾

Flat Profile

Search: Search Query (No Grouping)

{main}

Incl.	Self	Called	Function	Loc
100.00	63.50	(0)	{main}	0-n
35.60	35.60	10 000 001	php:fgetcsv	php
0.90	0.90	1	php:fopen	php
0.00	0.00	1	php:ksort	php

Types Callers All Callers Callee Map Source Code

#	Time_(10ns)	Time_(10ns)	Source
0			--- From '/home/bohuslavsimek/Public/TEST-1brc/0-naive.php' ---
1	63.50	63.50	<?php
2			\$stations = [];
3			\$fp = fopen('measurements-medium3.txt', 'r');
4	0.90	0.90	■ 1 call to 'php:fopen' (php:internal)
5			while (\$data = fgetcsv(\$fp, null, ',')) {
6	35.60	35.60	■ 10000001 calls to 'php:fgetcsv' (php:internal)
7			if (!isset(\$stations[\$data[0]])) {
8			\$stations[\$data[0]] = [
9			\$data[1],
10			...
11			}
12			...
13			}
14			ksort(\$stations);
15		0.00	■ 1 call to 'php:ksort' (php:internal)
16			echo '{';
17			foreach(\$stations as \$k=>&\$station) {

Samples: 23K of event 'cycles:P', Event count (approx.): 96664162627

Children	Self	Command	Shared Object	Symbol
+ 35,23%	18,12%	php	php	[.] zend strtod
+ 22,28%	0,00%	php	[unknown]	[.] 0000000000000000
+ 19,17%	12,62%	php	php	[.] _is_numeric_string_ex.part.0
+ 13,74%	4,78%	php	php	[.] zend_hash_find
+ 10,75%	0,00%	php	php	[.] zend_execute
+ 9,68%	0,00%	php	[unknown]	[.] 0x0000001600000001
+ 8,97%	5,02%	php	php	[.] zend_smart_strcmp
+ 7,45%	4,49%	php	php	[.] zif_substr
+ 7,32%	2,35%	php	php	[.] _emalloc
+ 6,91%	0,00%	php	[unknown]	[.] 0xfa1e0ff327ff41ff
+ 5,98%	2,69%	php	php	[.] add_function_slow
+ 4,84%	1,87%	php	php	[.] zend_compare
+ 4,52%	2,28%	php	php	[.] zend_string_equal_val
+ 4,34%	2,97%	php	php	[.] _php_stream_get_line
+ 4,26%	0,00%	php	[unknown]	[.] 0xfffffa78000000d0
+ 4,15%	0,01%	php	[kernel.kallsyms]	[k] do_syscall_64
+ 4,14%	0,00%	php	[kernel.kallsyms]	[k] entry_SYSCALL_64_after_hwframe
+ 4,11%	0,01%	php	libc.so.6	[.] read
+ 4,10%	0,00%	php	[unknown]	[k] 0x0000002400000006
+ 3,99%	0,02%	php	[kernel.kallsyms]	[k] __x64_sys_read
+ 3,85%	1,08%	php	libc.so.6	[.] __memchr_avx2
+ 3,40%	2,48%	php	php	[.] zif_fgets
+ 3,21%	2,84%	php	php	[.] _efree
+ 3,20%	1,84%	php	php	[.] zif strpos
+ 2,85%	1,75%	php	libc.so.6	[.] memmove_avx_unaligned_erms
+ 2,83%	1,59%	php	php	[.] ZEND_ASSIGN_DIM_OP_SPEC_VAR_CONST_HANDLER
+ 2,73%	0,00%	php	[unknown]	[.] 0x4cf7014c0c776349
+ 2,68%	0,00%	php	[kernel.kallsyms]	[k] x64_sys_call
+ 2,63%	0,01%	php	[kernel.kallsyms]	[k] ksys_read
+ 2,63%	0,03%	php	[kernel.kallsyms]	[k] vfs_read
+ 2,41%	0,00%	php	[unknown]	[.] 0x0000000000000005
+ 2,39%	1,49%	php	php	[.] zend_hash_func
+ 2,37%	0,08%	php	[kernel.kallsyms]	[k] filemap_read
+ 2,37%	0,01%	php	[kernel.kallsyms]	[k] ext4_file_read_iter
+ 2,36%	1,52%	php	opcache.so	[.] zend jit_symtable_lookup_rw
+ 2,35%	0,01%	php	[kernel.kallsyms]	[k] generic_file_read_iter
+ 2,35%	1,16%	php	php	[.] _zend_try_convert_scalar_to_number.isra.0
+ 2,33%	0,76%	php	php	[.] _erealloc

Perf - slowest functions

1. ~~php_fgetcsv_lookup_trailing_spaces~~
2. zend strtod
3. ~~php_fgetcsv~~
4. is_numeric_string_ex.part.0
5. zend_execute
6. **zend_hash_find**

How can I find out if this is connected to something I can change?

- A. Hash tables are used in PHP arrays.
- B. I could Google it.
- C. I could ask AI.
- D. Sacrifice a rubber duck to the code gods.

How can I find out if this is connected to something I can change?

- Hash tables are used in PHP arrays.



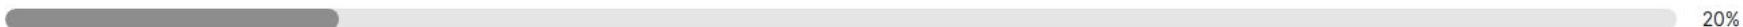
- I could Google it.



- I could ask AI.



- Sacrifice a rubber duck to the code gods.



`zend_hash_find?`
`Hash table?`

**Can I, as a PHP programmer,
limit the number of calls to the
zend_hash_find function?**

Can I, as a PHP programmer, limit the number of calls to the `zend_hash_find` function?



Yes, as a PHP programmer, you can indirectly limit the number of calls to `zend_hash_find` by optimizing how you interact with PHP arrays and hash tables. The `zend_hash_find` function is a low-level C function in the Zend Engine, which PHP uses to look up values in arrays (hash tables). While you cannot directly modify the engine to limit calls to `zend_hash_find`, you can optimize your PHP code to minimize unnecessary lookups. Here are some strategies:

1. Cache Results of Expensive Lookups

If you are performing multiple lookups on the same key, store the result in a variable instead of repeatedly querying the array:

```
php Copy code

// Inefficient
for ($i = 0; $i < 100; $i++) {
    $value = $array['key']; // zend_hash_find is called every time
    // Process $value
}

// Efficient
$value = $array['key']; // zend_hash_find called once
for ($i = 0; $i < 100; $i++) {
    // Process $value
```

```
12 } else {
13     $stations[$data[0]][3]++;
14     $stations[$data[0]][2] += $data[1];
15     if ($data[1] < $stations[$data[0]][0]) {
16         $stations[$data[0]][0] = $data[1];
17     }
18     if ($data[1] > $stations[$data[0]][1]) {
19         $stations[$data[0]][1] = $data[1];
20     }
21 }
22 } // while Loop end
```

Lets try it!

Use references

```
-     $stations[$city][3]++;
-     $stations[$city][2] += $temp;
-     if ($temp < $stations[$city][0]) {
-         $stations[$city][0] = $temp;
```

```
+     $station = &$stations[$city];
+     $station[3]++;
+     $station[2] += $temp;
+     if ($temp < $station[0]) {
+         $station[0] = $temp;
```

Use references

**Before
294.80 sec**



**After
276.06 sec**

6.36% faster!

Samples: 22K of event 'cycles:P', Event count (approx.): 93830033974

Children	Self	Command	Shared Object	Symbol
+ 39,71%	18,94%	php	php	[.] zend strtod
+ 23,89%	0,00%	php	[unknown]	[k] 0000000000000000
+ 21,59%	13,03%	php	php	[.] _is_numeric_string_ex.part.0
+ 10,75%	0,00%	php	php	[.] zend_execute
+ 9,32%	5,33%	php	php	[.] zendi_smart_strcmp
+ 7,69%	4,77%	php	php	[.] zif_substr
+ 7,30%	2,45%	php	php	[.] emalloc
+ 6,37%	0,00%	php	[unknown]	[.] 0x0000001600000001
+ 5,61%	0,00%	php	[unknown]	[.] 0xfale0ff327ff41ff
+ 5,14%	2,10%	php	php	[.] zend_compare
+ 5,09%	3,35%	php	php	[.] _php_stream_get_line
+ 4,94%	0,00%	php	[unknown]	[.] 0x0000001a00000001
+ 4,32%	1,14%	php	libc.so.6	[.] __memchr_avx2
+ 3,96%	2,17%	php	php	[.] ZEND_ASSIGN_REF_SPEC_CV_VAR_HANDLER
+ 3,94%	1,36%	php	php	[.] zend_hash_find
+ 3,85%	2,38%	php	php	[.] add_function_slow
+ 3,60%	2,10%	php	php	[.] zif strpos
+ 3,50%	2,62%	php	php	[.] _efree
+ 3,19%	0,00%	php	[unknown]	[.] 0xfffffa78000000b0
+ 3,17%	2,16%	php	php	[.] zif fgets
+ 2,98%	1,71%	php	php	[.] zend_hash_func
+ 2,92%	1,88%	php	libc.so.6	[.] __memmove_avx_unaligned_erms
+ 2,77%	0,01%	php	[kernel.kallsyms]	[k] entry_SYSCALL_64_after_hwframe
+ 2,76%	0,02%	php	libc.so.6	[.] read
+ 2,75%	0,00%	php	[kernel.kallsyms]	[k] do_syscall_64
+ 2,74%	0,00%	php	[unknown]	[.] 0x0000002400000006
+ 2,72%	0,00%	php	[unknown]	[.] 0x0f03087b83f3014c
+ 2,72%	0,00%	php	[unknown]	[.] 0xfffffa00000000d0
+ 2,70%	0,00%	php	[unknown]	[.] 0x0000704a43c141b0
+ 2,68%	1,47%	php	php	[.] ZEND_ASSIGN_DIM_OP_SPEC_CV_CONST_HANDLER
+ 2,56%	0,02%	php	[kernel.kallsyms]	[k] __x64_sys_read
+ 2,51%	1,40%	php	php	[.] zend_try_convert_scalar_to_number.isra.0

Perf - slowest functions

1. ~~php_fgetcsv_lookup_trailing_spaces~~
2. zend strtod
3. ~~php_fgetcsv~~
4. is_numeric_string_ex.part.0
5. zend execute
6. ~~zend_hash_find~~

TL;DR
Minimise
expansive code
calling

TL;DR
"try it and see"

**Check for
redundancy...**

```
12 } else {
13     $station = &$stations[$city];
14     $station[3]++;
15     $station[2] += $temp;
16     if ($temp < $station[0]) {
17         $station[0] = $temp;
18     }
19     if ($temp > $station[1]) {
20         $station[1] = $temp;
21     }
22 }
23 } // while Loop end
```

Check for Redundancy

- Temperature can't exceed max or fall below min simultaneously.
- Both conditions are mutually exclusive.

Remove extra comparison

```
$station[2] += $temp;  
if ($temp < $station[0]) {  
    $station[0] = $temp;  
}  
- if ($temp > $station[1]) {  
+ } else if ($temp > $station[1]) {  
    $station[1] = $temp;  
}
```

Remove extra comparison

Before
276.06 sec



After
250.58 sec

9.23% faster!

Perf - slowest functions

1. ~~php_fgetcsv_lookup_trailing_spaces~~
2. zend strtod
3. ~~php_fgetcsv~~
4. is_numeric_string_ex.part.0
5. zend_execute
6. ~~zend_hash_find~~

zend strtod and is_numeric_string_ex.part.0



In PHP's Zend Engine, `zend strtod` and `is_numeric_string_ex.part.0` are key parts of its handling of numeric values and string conversion. Here's a breakdown:

zend strtod

- **Purpose:** This function is responsible for converting a string to a floating-point number (`double`) in a locale-independent way.
- **Key Use Case:** It's utilized internally by PHP when a string needs to be parsed as a floating-point number. For example:

php

Copy code

```
$number = (float)"3.14"; // zend strtod likely plays a role here.
```

- **Implementation Notes:**

- It bypasses locale-specific behavior to ensure consistency regardless of system settings.
- This is crucial for avoiding issues where, for example, a locale might use commas (`1,23`) instead of dots (`1.23`) for decimals.

is_numeric_string_ex.part.0

- **Purpose:** This is an internal function used to determine whether a string can be interpreted as a numeric value (either an integer or a float). It may also indicate how the string should be interpreted.

Dynamic Typing

- Variables in PHP do not have fixed types.
- A variable's type can change at runtime.
- Is often dynamic - `is_numeric_string_ex`
- Can be expansive.

add_function_slow



+ 1,28% 0,77% php [.] add_function_slow

```
static zend_never_inline zend_result ZEND_FASTCALL add_function_slow(zval *result, zval *op1, zval *op2)
{
    // parts omitted for clarity
    zval op1_copy, op2_copy;
    if (UNEXPECTED(zendi_try_convert_scalar_to_number(op1, &op1_copy) == FAILURE)
        || UNEXPECTED(zendi_try_convert_scalar_to_number(op2, &op2_copy) == FAILURE)) {
        zend_binop_error("+", op1, op2);
        if (result != op1) {
            ZVAL_UNDEF(result);
        }
        return FAILURE;
    }

    if (add_function_fast(result, &op1_copy, &op2_copy) == SUCCESS) {
        return SUCCESS;
    }
    // parts omitted for clarity
}
```

How to limit impact of dynamic types in this case?

- A. Typecast to float ASAP, hope for PHP magic.
- B. Strictly validate input types before processing.
- C. Use declare(strict_types=1).
- D. Threaten PHP with a Java rewrite.

How to limit impact of dynamic types in this case?

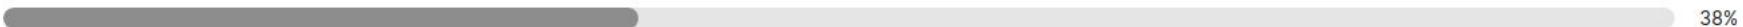
- Typecast to float ASAP, hope for PHP magic.



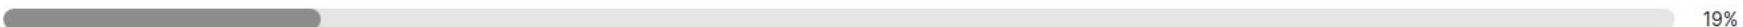
- Strictly validate input types before processing.



- Use declare(strict_types=1).



- Threaten PHP with a Java rewrite.



Typecasting

```
while ($data = fgets($fp, null)) {
    $pos = strpos($data, ';');
    $city = substr($data, 0, $pos);
-    $temp = substr($data, $pos+1, -1);
+    $temp = (float) substr($data, $pos+1, -1);
    if (!isset($stations[$city])) {
        $stations[$city] = [
            $temp,
```

Typecasting

**Before
250.58 sec**



**After
159.45 sec**

36.37% faster!

Samples: 12K of event 'cycles:P', Event count (approx.): 52597824389				
Children	Self	Command	Shared Object	Symbol
+	23,08%	0,00%	php	[unknown]
+	17,84%	0,00%	php	zend_execute
+	17,81%	9,51%	php	zend strtod
+	15,39%	0,00%	php	0xfa1e0ff327ff41ff
+	12,73%	7,98%	php	zif_substr
+	12,00%	4,21%	php	emalloc
+	9,23%	0,00%	php	0xfffffa58000000b0
+	7,75%	2,30%	php	zend_hash_find
+	7,68%	4,96%	php	php_stream_get_line
+	7,04%	3,64%	php	ZEND_ASSIGN_REF_SPEC_CV_VAR_HANDLER
+	6,88%	4,71%	php	_efree
+	5,87%	1,68%	php	_memchr_avx2
+	5,74%	3,76%	php	zif_fgets
+	5,59%	3,02%	php	zend_hash_func
+	5,46%	2,97%	php	zif strpos
+	5,23%	3,16%	php	_memmove_avx_unaligned_erm
+	4,68%	2,60%	php	ZEND_ASSIGN_DIM_OP_SPEC_CV_CONST_HANDLER
+	4,41%	1,75%	php	zend_hash_lookup
+	4,39%	0,00%	php	0x0000024000000006
+	4,39%	0,04%	php	[k] entry_SYSCALL_64_after_hwframe
+	4,38%	0,03%	php	[k] do_syscall_64
+	4,38%	0,02%	php	read
+	3,96%	0,00%	php	0xfffffa88000000b0
+	3,90%	1,90%	php	zend_string_equal_val
+	3,68%	0,03%	php	[k] __x64_sys_read
+	3,53%	2,58%	php	ZEND_CAST_SPEC_VAR_HANDLER
+	2,78%	0,92%	php	ZEND_PRE_INC_SPEC_VAR_RETVAL_UNUSED_HANDLER
+	2,62%	1,10%	php	_erealloc
+	2,49%	0,00%	php	0xfffffffff0000010a
+	2,45%	0,02%	php	[k] x64_sys_call
+	2,45%	0,00%	php	0x0f03087b83f3014c
+	2,45%	0,00%	php	0xfffff9c0000000d0
+	2,40%	1,70%	php	php_stream_locate_eol
+	2,39%	1,21%	php	zend_fetch_dimension_address_inner_RW_CONST
+	2,21%	0,00%	php	0x000001a000000001
+	2,18%	0,98%	php	_zend_mm_alloc
+	1,97%	0,09%	php	[k] vfs_read
+	1,97%	0,00%	[kernel.kallsyms]	

Samples: 12K of event 'cycles:P', Event count (approx.): 52597824389

Children	Self	Command	Shared Object	Symbol
- 23,08%	0,00%	php	[unknown]	[.] 0000000000000000
- 0				
- 17,84% zend_execute				
2,09% ZEND_ASSIGN_DIM_OP_SPEC_CV_CONST_HANDLER				
1,87% ZEND_PRE_INC_SPEC_VAR_RETVAL_UNUSED_HANDLER				
1,74% 0x58510ce012ea				
0,83% 0x58510ce011e5				
0,80% 0x58510ce012ec				
0,70% add_function				
0,69% 0x58510ce01370				
0,56% zend_is_true				
0,55% 0x58510ce012ac				
2,59% zend_hash_func				
- 1,55% 0x2820776100000106				
0,87% zval_get_double_func				

Perf - slowest functions

1. ~~php_fgetcsv_lookup_trailing_spaces~~
2. ~~zend_strtod~~
3. ~~php_fgetcsv~~
4. ~~is_numeric_string_ex.part.0~~
5. ~~z€ **CANCELED**ecute~~
6. ~~zend_hash_find~~

Use explode

```
while ($data = fgets($fp, null)) {  
- $pos = strpos($data, ';');  
- $city = substr($data, 0, $pos);  
- $temp = (float) substr($data, $pos+1, -1);  
+ [$city, $temperature] = explode(';', $data, 2);  
+ $temp = (float) $temperature;  
if (!isset($stations[$city])) {
```

Use explode

Before
159.45 sec



After
158.73 sec

0.45% “faster” :-/

What to do next?

- A. Rewrite portions in other language?
- B. Look how others solve it?
- C. Use parallelisation?
- D. Finally gave up!

What to do next?

- Rewrite portions in other language?



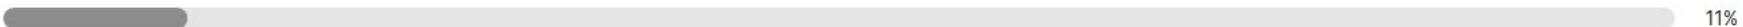
- Look how others solve it?



- Use parallelisation?



- Finally gave up!



Own extension in C - fastfileloader

```
$stations = [];
- $fp = fopen('measurements.txt', 'r');
- while ($data = fgets($fp, null)) {
+ $fp = fastfileloader_open(__DIR__.'/'.$data);
+ while ($data = fastfileloader_readline($fp)) {
    [$city, $temperature] = explode(';', $data, 2);
```

Naive simple implementation?



It's
slower

It's
slower

Own extension in C - fastfileloader

- It's 19.17% **slower** :-/
- Before 158.73 sec
- After **189.16** sec
- Reason: naive simple implementation.

Own extension in C - readline

```
PHP_FUNCTION(fastfileloader_readline) {
    // omitted for clarity

    if (fgets(buffer, sizeof(buffer), file) != NULL) {
        RETVAL_STRING(buffer);
    } else {
        if (feof(file)) {
            RETVAL_FALSE; // Return FALSE on end of file
        } else {
            php_error_docref(
                NULL, E_WARNING, "Error reading from file"
            );
            RETVAL_FALSE;
        }
    }
}
```

TL;dr - low level file access is need

- Need to redesign extension for peak perf.
- You need to know the OS and system things.
- In Linux mmap

Tl;dr - low level file access is need

```
PHP_FUNCTION(fastfileloader_open) {
    // removed for sake of clarity

    data = mmap(0, fileInfo.st_size, PROT_READ, MAP_PRIVATE, fd, 0);

    // removed for sake of clarity

    RETURN_RES(zend_register_resource(obj, le_file_resource));
}

PHP_FUNCTION(fastfileloader_readline) {
    // removed for sake of clarity

    char *end = memchr(start, '\n', obj->size - offset);

    // removed for sake of clarity

    RETVAL_STRINGL(start, len);
}
```

Own extension in C - second try

- Virtually same (speed up 0.77%) :-/
- Before 158.73 sec
- After 158.50 sec
- This is no worth it...

Own extension in C - second try

Before
158.73 sec



After
158.50 sec

0.77% “faster” :-/

TL;DR

**You are not that
smart.**

Inspiration from others

Parallelisation beat them all...

- Fastest 1BRC solutions uses parallelisation.
- In PHP extension **parallel**.
- Thread support in constraints of CSP model.

Last resort...



Start threads

```
$chunks = get_file_chunks($file, $threads_cnt); // $chunks = [[0, 123], [124, 256]]  
  
$futures = [];  
  
for ($i = 0; $i < $threads_cnt; $i++) {  
    $runtime = new \parallel\Runtime();  
    $futures[$i] = $runtime->run( // $chunks =  
        $process_chunk,  
        [  
            $file,  
            $chunks[$i][0],  
            $chunks[$i][1]  
        ]  
    );  
}
```

Thread code

```
$process_chunk = function (string $file, int $chunk_start, int $chunk_end): array {
    $stations = [];
    $fp = fopen($file, 'rb');
    fseek($fp, $chunk_start);
    while ($data = fgets($fp)) {
        $chunk_start += strlen($data);
        if ($chunk_start > $chunk_end) {
            break;
        }
        // regular processing
    }
    return $stations;
};
```

Synchronise and merge result

```
$results = [];

for ($i = 0; $i < $threads_cnt; $i++) {
    $chunk_result = $futures[$i]->value(); // blocks until a result is available,
    foreach ($chunk_result as $city => $measurement) {
        if (isset($results[$city])) {
            $result = &$results[$city];
            $result[2] += $measurement[2];
            $result[3] += $measurement[3];
            if ($measurement[0] < $result[0]) {
                $result[0] = $measurement[0];
            } else if ($measurement[1] > $result[1]) {
                $result[1] = $measurement[1];
            }
        } else {
            $results[$city] = $measurement;
        }
    }
}
```

Print result

```
ksort($results);

echo '{', PHP_EOL;
foreach($results as $k=>&$station) {
    echo "\t", $k, '=', $station[0], '/',
    number_format($station[2]/$station[3], 1),
    '/', $station[1], ',', PHP_EOL;
}
echo '}', PHP_EOL;
```

Parallelisation beat them all...

**Before
158.50 sec**



**After
29.81 sec**

81.19% faster!

Parallelisation beat them all...

- It's 81.19% faster!
- Before 158.50 sec
- With 20 threads 29.81 sec
- But you need extension and ZTS PHP build.

TL;DR

**Take inspiration
from others.**

Boss seems happy!



Skip directly to parallelisation?

- fgets (2nd) version as a base version
- It's 116.54% slower!
- 64.55 sec vs 29.81 sec
- macro-optimization/architectural change

Time to address the elephant in the room...



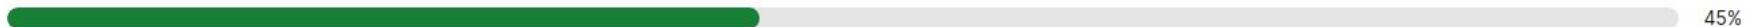
How does PHP compare with others?

Who will be faster Java or PHP?

- A. Java
- B. PHP
- C. I gave up, show the results!

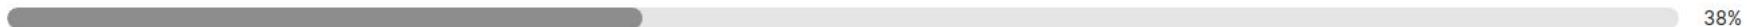
Who will be faster Java or PHP?

Java



45%

PHP



38%

I gave up, show the results!



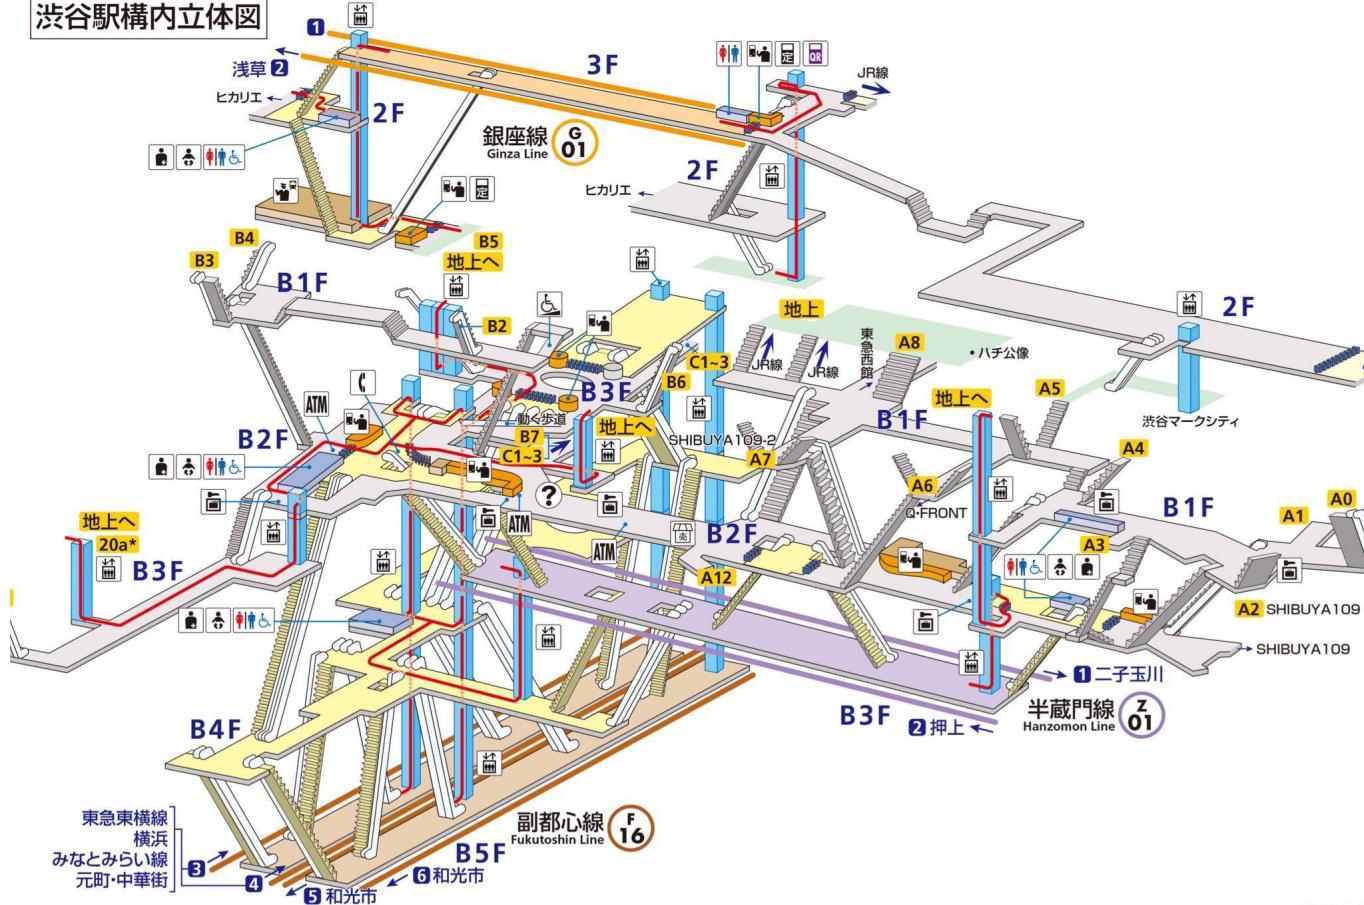
17%

In broad context...

- Java
 - Baseline - 131.85 sec
 - Fastest - 5.2 sec
- PHP - 29.81 sec
- C - 5.25 sec

**How best solution
in JAVA looks?**

渋谷駅構内立体図



[2024.8]

```
1 private static long scanNumber(Scanner scanPtr) {
2     long numberWord = scanPtr.getLongAt(scanPtr.pos() + 1);
3     int decimalSepPos = Long.numberOfTrailingZeros(~numberWord & 0x10101000L);
4     long number = convertIntoNumber(decimalSepPos, numberWord);
5     scanPtr.add((decimalSepPos >>> 3) + 4);
6     return number;
7 }
8
9 private static long convertIntoNumber(int decimalSepPos, long numberWord) {
10     int shift = 28 - decimalSepPos;
11     long signed = (~numberWord << 59) >> 63;
12     long designMask = ~(signed & 0xFF);
13     long digits = ((numberWord & designMask) << shift) & 0x0F000F0F00L;
14     long absValue = ((digits * 0x640a0001) >>> 32) & 0x3FF;
15     return (absValue ^ signed) - signed;
16 }
```

Why is Java version so fast?

- Developed by lead of GraalVM.
- Native AOT binary.
- Parallelize I/O
- Uses integers for calculation.

Why is Java version so fast?

- Custom hash table.
- Uses sun.misc.Unsafe.
- Branchless conversion between formats.
- Branch prediction.

**My pragmatic
solution?**

Personal choice, DuckDB, ~33.96 sec

```
WITH src AS (
    SELECT station_name,
    MIN(measurement) AS min_measurement,
    CAST(AVG(measurement) AS DECIMAL(8,1)) AS mean_measurement,
    MAX(measurement) AS max_measurement
    FROM READ_CSV(
        'measurements.txt', header=false, columns= {'station_name':'VARCHAR','measurement':'double'}, delim=';', parallel=true
    )
    GROUP BY station_name
)
SELECT '{}' ||
    ARRAY_TO_STRING(
        LIST_SORT(
            LIST(station_name || '=' || CONCAT_WS('/',min_measurement, mean_measurement, max_measurement))
        ), ','
    )
    ) || '}' AS "1BRC"
FROM src;
```

Recapitulation

Recapitulation

- Check configuration and runtime.
- Get to know your tools.
- Find out what is slow.

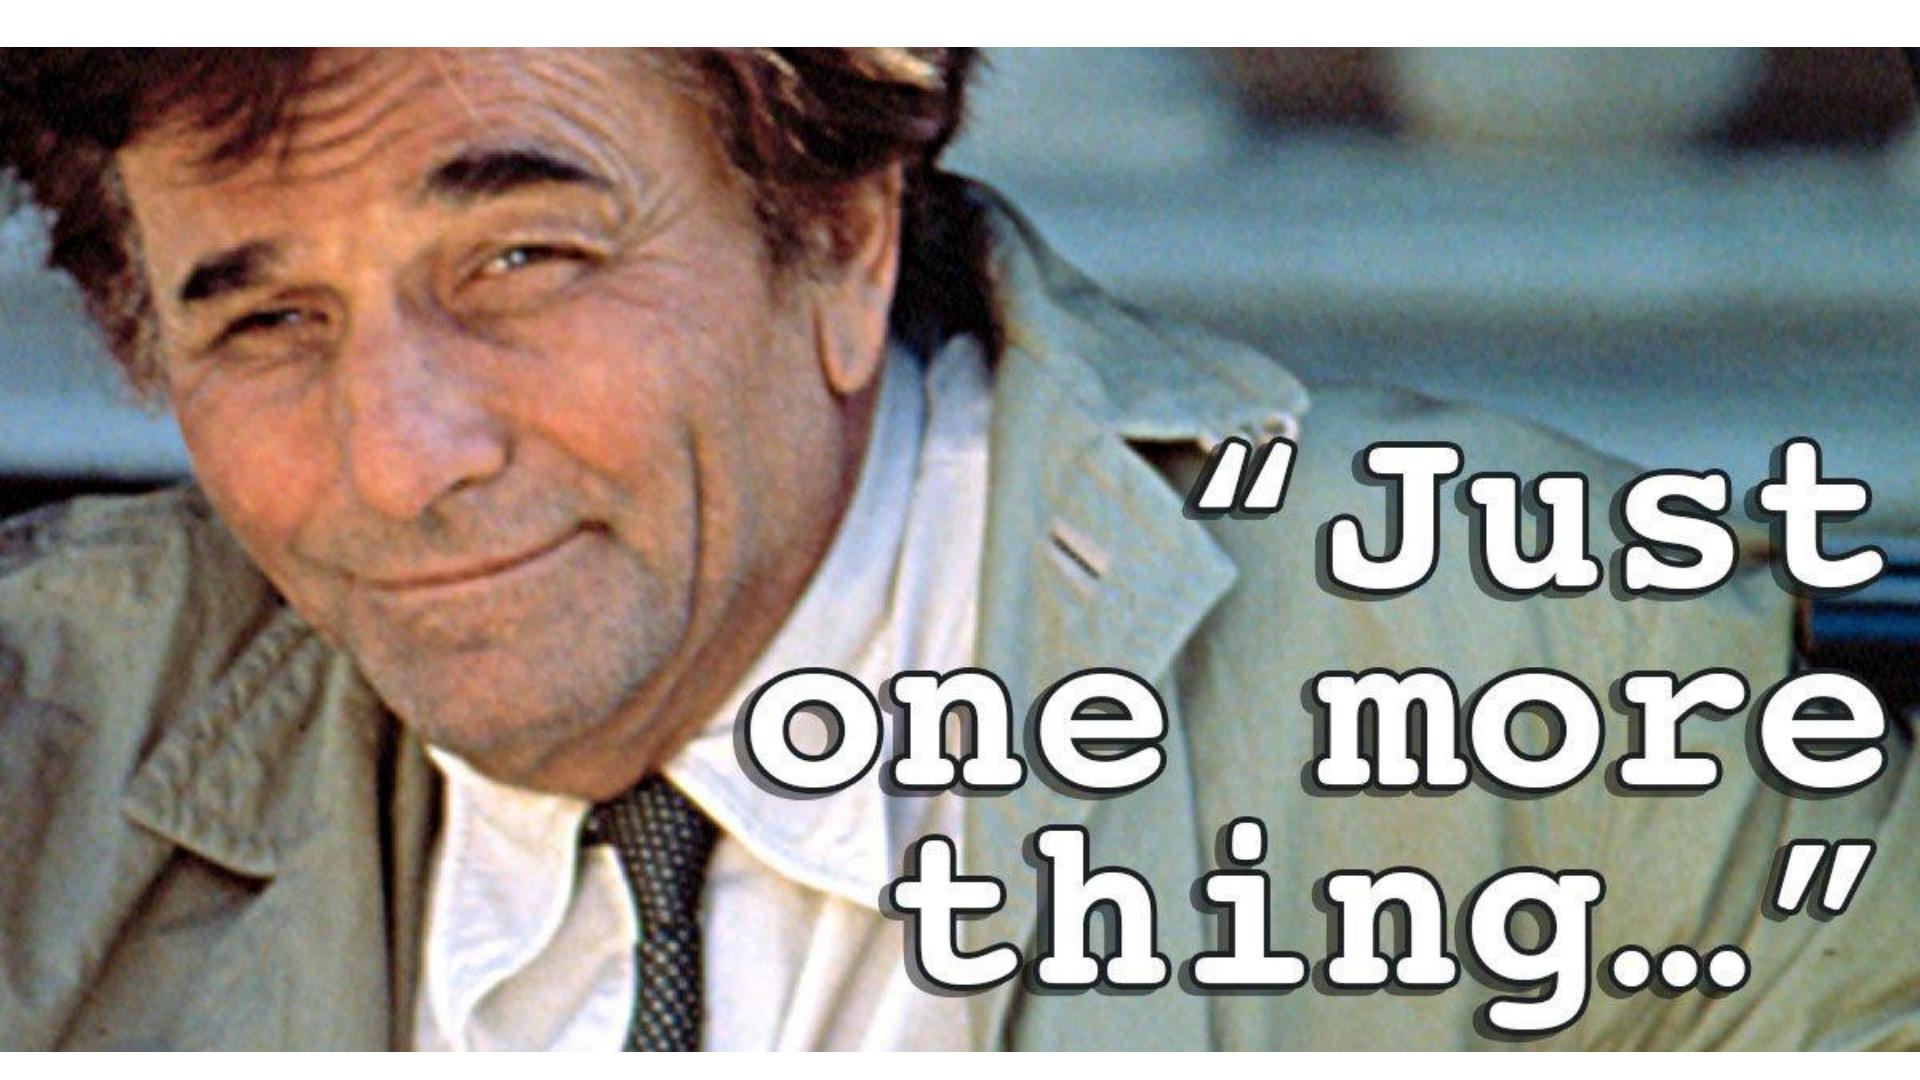
Recapitulation

- Find out what's essential.
- Do only what's essential.
- "Try it and see" – great for AI.

Recapitulation

- Minimise expansive code calling.
- You are not that smart.
- Take inspiration from others.

Was it worth it?

A close-up photograph of actor Robert De Niro's face. He has a mischievous, knowing smile and is looking slightly off-camera to his left. He is wearing a light-colored, possibly beige or tan, jacket over a white collared shirt and a dark, patterned tie. The background is blurred, showing what appears to be a stone wall or doorway.

**"Just
one more
thing..."**

Quiz result

Quiz result

1	Vojta	7/9	4:03
2	Karel	5/9	4:10
3	Maňas	5/9	4:48
4	Your Boss	5/9	4:58
5	Šáník	5/9	5:33
6	Jiří	4/9	3:14
7	Miroslav Hatina	4/9	3:39
8	Perry	4/9	3:58
9	Marx	4/9	4:02
10	Jiří PPH	4/9	4:37
11	Prochi	4/9	4:38

“That’s all Folks!”

Isberg®

Moderní PHP: Využijte sílu nových funkcí!

- **Co je nového v PHP?** – Představíme klíčové novinky z posledních verzí PHP, rychlostní testy, moderní syntax a nové funkcionality
- **Tipy na moderní kód** – Praktické ukázky, jak psát čitelnější, efektivnější a udržitelnější kód
- **Jak na migraci na nové verze?** – Rady a postupy, jak efektivně přejít na nejnovější verze PHP ve vašich projektech



Osteria Garage
Salonek v 1. patře
Částkova 56, Plzeň



29. 01. 2025
18:00 hod

